

RepTrap: A Novel Attack on Feedback-based Reputation Systems *

Yafei Yang[†], Qinyuan Feng*, Yan Lindsay Sun[†] and Yafei Dai*

[†] University of Rhode Island, Kingston, RI, USA

Email: {yafei, yansun}@ele.uri.edu

*CNDS Lab, Peking University, Beijing, China

Email: {fqy, dyf}@net.pku.edu.cn

ABSTRACT

Reputation systems are playing critical roles in securing today's distributed computing and communication systems. Similar to other security mechanisms, reputation systems can be under attack. In this paper, we report the discovery of a new attack, named RepTrap(Reputation Trap), against feedback-based reputation systems, such as those used in P2P file-sharing systems and E-commerce websites(e.g. Amazon.com). We conduct an in-depth investigation on this new attack, including analysis, case study, and performance evaluation based on real data and realistic user behavior models. We discover that the RepTrap is a strong and destructive attack that can manipulate the reputation scores of users, objects, and even undermine the entire reputation system. Compared with other known attacks that achieve the similar goals, the RepTrap requires less effort from the attackers and causes multi-dimensional damage to the reputation systems.

Categories and Subject Descriptors

C.2.4 [Computer-Communication Networks]: Distributed Systems; C.2.0 [Computer-Communication Networks]: General—Security and Trust

General Terms

Security, Attacks, Algorithms

Keywords

Reputation System, Feedback, Trust

1. INTRODUCTION

Word-of-mouth, one of the most ancient mechanisms in the history of human society, is gaining new significance in the Internet [7, 16]. The *online reputation systems*, also known as the online feedback mechanisms, are creating large scale, virtual word-of-mouth networks in which individuals share opinions and experiences on a wide range of topics, including products, companies, digital content and even other people. For

*This work is partial supported by NSF(award 0643532), NSFC(award 60673183), and Doctoral Funding of MOE (20060001044)

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

SecureComm '08 September 22 - 25, 2008, Istanbul, Turkey
Copyright 2008 ACM 978-1-60558-241-2 ...\$5.00.

example, Epinions.com encourages Internet users to rate practically any kind of businesses. Citysearch.com solicits and displays user feedback on restaurants, bars, and performances. YouTube.com recommends video clips based on viewers' feedback and some P2P file-sharing systems[24] do the same for the files shared by the users.

The reputation systems play significant roles in (1) assisting users' decision-making, (2) encouraging trustworthy behavior, and (3) deterring participation by unskilled or dishonest users. Reputation systems are having increasing influence on purchasing decision of consumers, online digital content distribution, and even political campaigns [7].

Meanwhile, the *manipulation* of such systems is rapidly growing. Firms post biased ratings and reviews to praise their own products or bad-mouth the products of their competitors. Political campaigns promote positive video clips and hide negative video clips by inserting unfair ratings at YouTube.com. There is ample evidence that such manipulation takes place [8]. In February 2004, due to a software error, Amazon.com's Canadian site mistakenly revealed the true identities of some book reviewers. It turned out that a sizable proportion of those reviews were written by the books' own publishers, authors, and competitors [12]. The scammers are creating sophisticated programs that mimic legitimate YouTube traffic and provide automated feedback for videos and other content they wish to promote[13]. Some eBay users are artificially boosting their reputation by buying and selling feedbacks[3]. Collaborative manipulation of feedback based reputation systems is a growing threat. This threat is hurting consumers and will also hurt the business hosting such systems [4].

In the current literature, the research on attacks against reputation systems is still immature. The existing threat models are rather straightforward[14, 22]. The known attacks can be classified according to their attack goals: *self-promoting* (i.e. falsely increasing reputation), *slandering* (i.e. falsely reducing reputation), *whitewashing* (i.e. repairing reputation after bad behaviors), and *denial-of-service* (i.e. making the system wrongly accuse honest users or objects). Attackers achieve the above goals through a variety of methods[14]. Each attacker can acquire multiple identities through a sybil attack[9, 29, 28]. All the identities under the control of the attackers can (1) provide positive feedbacks for self-promoting, (2) provide negative feedbacks for slandering; (3) behave honestly to the objects that the attackers are not interested in or simply register as a new user for whitewashing[19, 20]; and (4) subvert the underlying detection algorithms to make honest users/objects look suspicious[22].

Whereas current research primarily focuses on individual attack strategies, the intelligent human attackers can create sophisticated attacks by integrating different manipulation strategies. In this paper, we report the discovery of a new attack, named *RepTrap*, against feedback-based reputation systems, and conduct an in-depth investigation on this new attack.

RepTrap is applicable to reputation systems that calculate

key reputation scores based on *feedbacks* as well as metrics describing whether users provide honest feedbacks. These metrics are often referred to as *feedback reputation*. Although RepTrap is applicable to a broad range of systems, we choose P2P file-sharing system as the evaluation platform for three reasons. First, there are real trace data[18] and mature models describing user behaviors in P2P networks[11, 26]. Second, P2P networks, such as Kazaa [2], eMule [1] and Maze[18], are becoming the most important infrastructure to share content and services among users in a distributed way. Third, reputation systems play critical roles in P2P file-sharing systems as a method to filter out fake content (pollution) and malicious users [19, 10]. The simulation results have demonstrated that many current reputation systems are highly vulnerable to this new attack. With the same attack effort, RepTrap can significantly increase the attackers' chance of success. Furthermore, this attack can undermine the users' incentive to participate and contribute to the application systems.

As a summary, our *contributions* are (1) discovery of a new and powerful attack against feedback-based reputation systems, which reveals a severe vulnerability in prevalent reputation systems; (2) formal attack strategies to maximize the effectiveness of this attack; and (3) an in-depth study on the impact of the attack in P2P file-sharing systems. The rest of the paper is organized as follows. An overview of feedback-based reputation systems and related work are presented in Section II. Section III presents the RepTrap attack, including basic ideas, case studies and detailed attack methods. Section IV demonstrates the performance evaluation. Some additional discussions are provided in Section V, followed by the conclusion in Section VI.

2. REPUTATION SYSTEM REVIEW AND RELATED WORK

Reputation systems have been used in many application scenarios. Instead of studying a specific reputation system for a specific application, we would like to broaden the scope of this investigation by taking the following 3 steps.

- Step 1: We summarize the common properties of prevalent application scenarios of reputation systems, and build architecture of reputation systems without specifying the algorithms for reputation calculation.
- Step 2: Based on this architecture, we discover the new attack that can undermine the foundation of reputation systems.
- Step 3: To demonstrate the performance of the attack, specific algorithms and application scenario are put back into the picture. We will quantitatively evaluate the attack in a P2P file-sharing system based on real trace data.

Step 1 and 2 will be described in this section. Step 3 will be discussed in Section 3 and demonstrated in Section 4.

2.1 Application Scenarios

Feedback-based reputation systems are used in many online applications as discussed in Section 1. These application systems have two common properties.

First, they have reputation mechanisms that provide information to users for assisting their decision-making. For example, P2P file-sharing systems can provide scores representing the quality of files such that users can decide which file to download[24]; product rating systems provide rating scores describing the quality of the products to guide users' purchasing decisions; eBay maintains the reputation scores of sellers/buyers such that one can decide whether to trust a seller/buyer. All above scores can be viewed as reputation, and are referred to

as the *primary reputation* in this work. In addition, the party, whose quality is described by the primary reputation score, is referred to as the *object*. An object in P2P file-sharing, product rating, and user rating system is a file, a product, and a user, respectively. In other words, as the first common property, these application systems *calculate and maintain primary reputation scores describing the quality of the objects*. So in the rest of the paper, the primary reputation is referred to as the *object quality reputation*.

Second, users can provide *feedbacks* about the quality of the objects, and the feedbacks are the primary source to determine the object quality reputation.

It is important to point out that many applications provide services in a *distributed* manner, such as P2P file-sharing systems in which users can directly interact with each other. However, the reputation system can be *centralized*. That is, there are one or several central authorities collecting feedbacks, publishing reputation scores, and detecting abnormalities. The central authority is referred to as Trust Manager (TM). It does not mean the system must have a central server. The central authority can be formed by a group of users. For example, even in a pure distributed P2P file-sharing network, we can build a reputation system with DHT[21]. In this paper, we focus on the usage of RepTrap in centralized reputation systems. Its extension to distributed reputation systems will be investigated in the future work.

As a short summary, the key concepts in the application scenarios that we care about in this paper are: *users, objects, feedback, and object quality reputation*.

2.2 Reputation System Architecture

Figure 1 shows the basic building blocks of the reputation systems and the relationship among these building blocks.

Object quality reputation: In most reputation systems, the primary goal is to determine the object quality reputation, which will assist the users to select high-quality objects, and assist the system to detect low-quality objects. As illustrated in Figure 1, the object quality reputation is determined by feedbacks from users, feedback reputation of users, and possibly other factors. The algorithm that calculates the object quality reputation from the above factors is referred to as the *QR Algorithm*.

Feedback and feedback reputation: Utilizing feedback has both advantages and disadvantages. As an advantage, feedbacks (if from honest users) can accurately reflect the quality of the objects and therefore allow a fast establishment of object quality reputation. On the other hand, the system must provide incentive to users such that they will provide honest feedbacks. More importantly, the usage of feedbacks leads to feedback-based attacks against reputation systems. In particular, malicious users can provide dishonest feedbacks and mislead the calculation of object quality reputation. More details will be discussed in Section 2.3.

The feedback reputation is calculated based on the users' previous feedback behaviors. The feedback behavior is often described by

- the number of honest feedbacks given by this user previously, denoted by G_{fb} ;
- the number of dishonest feedbacks given by this user previously, denoted by B_{fb} ;
- time when the feedbacks were given.

The algorithm that calculates user feedback reputation is referred to as the *FR Algorithm*.

Evidence collection: The Evidence Collection (EC) module provides inputs to the QR algorithm and the FR algorithm.

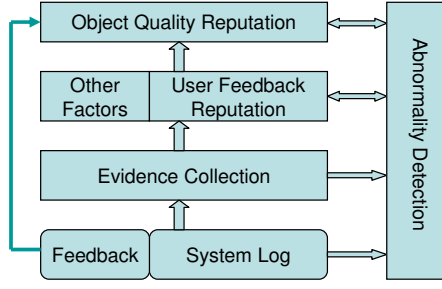


Figure 1: Architecture of Reputation System.

The module needs to judge whether a user’s feedback is honest and whether an object has high quality.

It is important to point out that the system usually does not know the ground truth about whether an object is high-quality or low-quality. Therefore, the EC often needs to estimate the ground truth. The common practice is to assume that the object quality reputation calculated by the reputation system is accurate. As we will discuss in the next section, this assumption introduces vulnerabilities into the reputation system.

Abnormal detection: This module detects low-quality objects, malicious users, attacks against application systems, attacks against reputation systems, and any other abnormalities.

Finally, we would like to point out that the current literature already provides ample choices for QR and FR algorithms. Examples include the beta-function based method in [15], the simple summation method in eBay [7], the entropy-based methods in [25]. An overview of more algorithms can be found in [16]. Meanwhile, the research on abnormality detection is still in its early stage.

2.3 Related Work

Generally speaking, the basic goal of malicious attacks against a reputation system is to boost or reduce the reputation of certain objects. There is an *arms race* between the attack and defense efforts.

Attack: To boost or reduce the object quality reputation, the most straightforward method is to insert dishonest feedbacks. This is referred to as the bad-mouthing attack[22]. (This attack has other names. For example, in eBay-like systems, ballot stuffing refers to a seller collecting positive feedbacks from faked transactions with friends; bad-mouthing refers to deliberately lowering others’ reputation through negative feedbacks[5, 6, 7].) In this paper, we adopt the broader definition of bad-mouthing in [14], and use **RepBad** to denote this attack.

Defense: To defeat the bad-mouthing attack, the reputation system adopts the feedback reputation [23]. Feedback reputation also has different names, such as “overall reputation score of users” in [17] and “trust in raters” in [27].

Attack: If feedback reputation is used, the attackers need to maintain their feedback reputation after providing dishonest feedbacks. The attackers can provide honest feedbacks to the objects that they are not interested in. This is referred to as *reputation recovery*[22] or *self-promoting* [14]. To avoid confusion, we use **RepSelf** to denote this attack in this paper.

Defense: To reduce the effectiveness of feedback reputation self-promoting, the system can limit the maximum number of feedbacks per user in a given time interval. This is reasonable because regular users do not use/evaluate a large number of objects in a short period of time. Therefore, the attackers must balance the usage of the limited number of feedbacks between attack and self-promoting, i.e. either to (1) provide a large number of dishonest feedbacks and have not-so-good feedback reputation or (2) to provide a small number of dishonest feedbacks and maintain a good feedback reputation. By limiting

the number of feedbacks per user, the reputation system can suppress either the *impact* of each dishonest feedback or the *number* of dishonest feedbacks.

3. REPTRAP: A NEW ATTACK AGAINST REPUTATION SYSTEMS

3.1 Basic Ideas

In self-promoting, the attackers only increase their own feedback reputation. How about reducing the honest users’ feedback reputation? Can the attackers hurt honest users’ feedback reputation and improve their own feedback reputation at the same time? In this paper, we move the attack-defense arms race one step further and propose a new attack strategy, called *Reputation Trap* (RepTrap).

RepTrap Attack: Attackers find some high quality objects that have a small number of feedbacks. Then, attackers provide a large number of negative feedbacks to these objects such that these objects are marked as low-quality by the system. That is, the system makes wrong judgement about the object quality. As a consequence, the system thinks the attackers’ negative feedbacks agree with the object quality and the feedbacks from honest users disagree with the object quality. Therefore, the feedback reputation of the attackers will be increased while the feedback reputation of honest users will be reduced.

RepTrap is applicable under three conditions. *First*, there exist high quality objects with a small number of feedbacks. Many existing applications, such as P2P file-sharing systems and Amazon.com, satisfy this condition. *Second*, the reputation system uses feedback reputation to calculate object quality reputation. *Third*, the system does not know the ground truth about the object quality and has to estimate the ground truth based on the object quality reputation. This is true in many large scale open systems. It is also noted that this condition may not be satisfied in some semi-centralized systems where a large number of experts are hired to evaluate product quality and the expert opinions override users’ feedbacks.

As a summary, *the RepTrap attack hurts the feedback reputation of honest users and boosts the feedback reputation of attackers by undermining the system’s estimation of object quality*. The RepTrap attack also has other effects, such as hurting good users’ incentive to collaboration, which will be discussed in Section 5.

3.2 Case Studies

To illustrate the basic idea of RepTrap, we create an example in a simplified reputation system. In this section, we first describe this simple reputation system, then present the example, and finally compare the effectiveness of various attacks in this example scenario.

3.2.1 A Simple Reputation System

This simple reputation system takes binary feedbacks. Positive feedback indicates high quality and negative feedback indicates low quality. One user can provide at most one feedback to one object. The beta-function based trust model is adopted to calculate reputation scores in the QR algorithm and FR algorithm. The EC module uses thresholding method. Particularly,

1. The feedback reputation of user u_j is calculated as $R^{fr}(u_j) = \frac{G_{fb}(u_j)+1}{G_{fb}(u_j)+B_{fb}(u_j)+2}$, where G_{fb} and B_{fb} are defined in Section 2.2.
2. The quality reputation of object O_k is calculated as a weighted average of feedback values with the weight factors as users’ feedback reputation. In particular,




Figure 2: Object-User Graph.

$$R^{qr}(O_k) = \frac{\sum_{u_j \in U_{pos}^k} R^{fr}(u_j)}{\sum_{u_j \in U_{pos}^k} R^{fr}(u_j) + \sum_{u_i \in U_{neg}^k} R^{fr}(u_i)}, \quad (1)$$

where U_{pos}^k and U_{neg}^k denote the set of users who give object O_k positive and negative feedbacks, respectively.

3. Update the G_{fb} and B_{fb} values using the following rules.
 - If $R^{qr}(O_k) \geq threshold$, O_k is marked as a high-quality object; if $R^{qr}(O_k) < threshold$, O_k is marked as a low-quality object.
 - If a user u_j provides a positive (or negative) feedback to an object marked with high (or low) quality, the value of $G_{fb}(u_j)$ is increased by 1. Otherwise, if the feedback and the object quality reputation do not match, $B_{fb}(u_j)$ is increased by 1.
4. If the values of G_{fb} and B_{fb} for any users are changed in step 3, go to step 1.

3.2.2 A Simple Application Scenario

We create a simple scenario with 5 objects: O_1, O_2, O_3, O_4 , and O_5 , and 3 honest users: u_1, u_2 , and u_3 . All the objects have high quality. u_1 gives four feedbacks; while u_2 and u_3 each gives two feedbacks.

$$\begin{aligned} O_1 \leftarrow u_1, O_2 \leftarrow u_1, O_3 \leftarrow u_1, O_4 \leftarrow u_1 \\ O_1 \leftarrow u_2, O_2 \leftarrow u_2; O_1 \leftarrow u_3, O_5 \leftarrow u_3 \end{aligned}$$

where $O_k \leftarrow u_j$ denotes that user u_j provides a feedback for object O_k . The relationship between objects and users is illustrated in Figure 2. To concisely describe the status of an object, we introduce the following notations.

$$R_{pos}^{fr}(O_k) = \sum_{u_j \in U_{pos}^k} R^{fr}(u_j), \quad R_{neg}^{fr}(O_k) = \sum_{u_i \in U_{neg}^k} R^{fr}(u_i).$$

Then, (1) becomes

$$R^{qr}(O_k) = \frac{R_{pos}^{fr}(O_k)}{R_{pos}^{fr}(O_k) + R_{neg}^{fr}(O_k)}. \quad (2)$$

For the example in this section, we set $threshold = 0.3$. And we can then calculate the reputations scores using the algorithms described in Section 3.2.1. When there are no malicious users, the feedback reputation of the honest users is:

$$u_1 : 0.83; \quad u_2 : 0.75; \quad u_3 : 0.75$$

and the object quality reputation is

$$O_1 : 1.00; \quad O_2 : 1.00; \quad O_3 : 1.00; \quad O_4 : 1.00; \quad O_5 : 1.00$$

Thus, all of the objects are marked as high quality.

3.2.3 Goal of Attacker

In this example, the **attacker's goal** is to make the system mark O_1 as a low-quality object while minimizing the *attack effort*.

The attack effort is described by two values (K, F) . Here, K is the number of user IDs under the attacker's control. These user IDs are also referred to as **malicious users**. F is the total number of feedbacks from the malicious users. Higher are the K and F values, more resources the attacker needs to spend. For an attacker, acquiring user IDs is usually harder than providing feedbacks, for two reasons. First, many systems limit the number of user IDs per IP address. Second, there are many techniques to defense against the Sybil attacks [29, 28] such that the cost for an attacker to control many user IDs increases greatly. Therefore, reducing the K value has higher priority than reducing the F value, from the attacker's point of view. Thus, the attacker first finds the minimal K value, and then finds the minimum F value given the minimal K value.

The attacker's goal is then translated into finding a way to provide feedbacks such that (1) the K value is minimized and (2) the F value is minimized given the minimal K value, under the constraint

$$R^{qr}(O_1) < threshold \quad (3)$$

3.2.4 Comparison among Different Attack Strategies

In this subsection, we compare RepTrap with RepBad and RepSelf, which are defined in Section 2.3. From (2) and (3), we derive the necessary and sufficient condition for making the reputation system mark an object (O_k) as low quality. The condition is:

$$R_{neg}^{fr}(O_k) > R_{pos}^{fr}(O_k) \cdot \frac{1 - threshold}{threshold} \quad (4)$$

In this case, if the attacker wants O_k to be marked as low-quality when there are no previous negative feedbacks given to O_k , the attacker needs to make some malicious users provide negative feedbacks to O_k and the summation of the malicious users' feedback reputation needs to be higher than $\frac{7}{3} R_{pos}^{fr}(O_k)$. We define the value of $\frac{7}{3} R_{pos}^{fr}(O_k)$ as the *hardness* value of an object. In this example, the hardness values of the objects are:

$$O_1 : 5.44; \quad O_2 : 3.69; \quad O_3 : 1.94; \quad O_4 : 1.94; \quad O_5 : 1.75$$

RepBad: In the RepBad attack, the feedback reputation for each malicious user is 0.5, and the hardness of O_1 is 5.44. Since $5.44/0.5 = 10.88$, the attacker needs at least 11 malicious users to successfully attack O_1 . In this case, the attack power is ($N = 11, K = 11$).

RepSelf: In the RepSelf attack, the malicious users first provide honest feedbacks to the objects that they do not want to attack. In this example, they can provide positive feedbacks to O_2, O_3, O_4 and O_5 , and accumulate their feedback reputation up to $(4 + 1)/(4 + 2) = 0.83$. Since $5.44/0.83 = 6.55$, at least 7 malicious users are needed to successfully attack O_1 .

Not all 7 malicious users need to provide 4 honest feedbacks. After enumerating all possible ways to perform self-promoting, we find that the malicious users need to provide at least 25 feedbacks. To achieve this,

- 3 malicious users provide positive feedbacks to 2 objects, and their feedback reputation becomes $(2 + 1)/(2 + 2) = 0.75$.
- 4 malicious users provide positive feedbacks to 3 objects, and their feedback reputation becomes $(3 + 1)/(3 + 2) = 0.8$.

Then, the summation of the feedback reputation of malicious users is $0.75 \times 3 + 0.8 \times 4 = 5.45 > 5.44$, which means they can

	Initial state	After 1 st round	After 2 nd round	After 3 rd round	After 4 th round
$R^{fr}(u_1)$	0.83	0.83	0.67	0.50	0.33
$R^{fr}(u_2)$	0.75	0.75	0.75	0.75	0.50
$R^{fr}(u_3)$	0.75	0.50	0.50	0.50	0.50
$R^{fr}(X_1)$	0.50	0.67	0.75	0.80	0.83
$R^{fr}(X_2)$	0.50	0.67	0.75	0.80	0.83
$R^{fr}(X_3)$	0.50	0.67	0.75	0.80	0.83
$R^{fr}(X_4)$	0.50	0.67	0.67	0.67	0.75
$\sum_{i=1}^4 R^{fr}(X_i)$	2.00	2.68	2.92	3.07	3.24
$Hardness(O_1)$	5.44	4.86	4.47	4.080	3.11
$Hardness(O_2)$	3.69	3.69	3.31	2.92	#
$Hardness(O_3)$	1.94	1.94	1.56	#	#
$Hardness(O_4)$	1.94	1.94	#	#	#
$Hardness(O_5)$	1.75	#	#	#	#

Table 1: Reputation and hardness values in case study

successfully attack O_1 . Thus, total 18 honest feedbacks and 7 dishonest feedbacks are given. In this case, the attack power is ($N = 7, K = 25$).

RepTrap: In the RepTrap attack, the malicious users provide dishonest feedbacks to unpopular high-quality objects. If a high-quality object is marked as low quality by the reputation system, this object is turned into a “**trap**”, which will mislead the feedback reputation of good users who give honest feedback to this trap.

It is difficult to find the optimal way to conduct the RepTrap attack. Instead, we just show one method to successfully attack O_1 using RepTrap. The attack power of this method is ($N = 4, K = 18$).

1. The initial feedback reputation of the four malicious users is 0.5. Since $0.5 \times 4 = 2$ is larger than the hardness of O_5 (i.e. 1.75), they can turn O_5 into a trap by providing negative feedbacks to O_5 . This has three consequences. First, the feedback reputation of the malicious users increases to $(1 + 1)/(1 + 2) = 0.67$. Second, the feedback reputation of good user u_3 is reduced to $(1 + 1)/(2 + 2) = 0.5$. Third, the value of $R_{pos}^{fr}(O_1)$ and the hardness of O_1 is reduced. The results are shown in the second column in Table 1, and we use X_1, X_2, X_3 and X_4 to denote the four malicious users.
2. After the first step, three malicious users can turn O_4 into a trap because the summation of their feedback reputation is larger than the hardness of O_4 , i.e. $0.67 \times 3 > 1.94$. Then, the feedback reputation of these three malicious users are increased to $(2 + 1)/(2 + 2) = 0.75$, and the feedback reputation of the honest user u_1 is reduced to $(3 + 1)/(4 + 2) = 0.67$. The hardness of O_1, O_2 and O_3 is also reduced, because their hardness depends on the feedback reputation of u_1 . See the third column in Table 1 for details.
3. The three malicious users in the second step can further turn O_3 into a trap because $0.75 \times 3 = 2.25$ is larger than the current hardness of O_3 . Then, the feedback reputation of these three malicious users are increased to $(3 + 1)/(3 + 2) = 0.8$, the feedback reputation of u_1 is reduced to $(2 + 1)/(4 + 2) = 0.5$, and the hardness of O_1 and O_2 continues to drop.
4. Similarly, all four malicious users can turn O_2 into a trap. Then, the feedback reputation of malicious users becomes 0.83 or 0.75, the feedback reputation of u_1 becomes 0.33, and the hardness of O_1 is reduced to 3.11.

5. Finally, the summation of feedback reputation of the four malicious users is $0.83 \times 3 + 0.75 = 3.24$, which is larger than 3.11. This means that the malicious users can now make the system mark O_1 as low quality. In total, the malicious users give $4 + 3 + 3 + 4 + 4 = 18$ feedbacks.

In this case study, RepTrap reduces the requirement on the number of malicious users by 64% when compared with RepBad, and by 43% when compared with RepSelf. RepTrap also requires 28% less feedbacks than RepSelf. In other words, RepTrap is a much stronger attack than the existing ones in manipulating the reputation system.

It is noted that malicious users provide many negative feedbacks in RepTrap. In systems such as eBay, this can be detected because positive feedbacks overwhelm negative feedbacks in such systems. However, in many systems, users give the negative feedbacks much more frequently than the positive feedbacks due to lack of incentive [7, 16]. In general, RepTrap cannot be effectively detected simply based on the number of negative feedbacks provided by individual user IDs.

3.3 Basic Concepts in RepTrap

As introduced in Section 3.1, the core idea of RepTrap is to degrade the feedback reputation of honest users by undermining the system’s estimation on object quality. In Section 3.2, a case study shows the effectiveness of RepTrap. To fully understand this new attack, we develop a *systematic procedure* to conduct the RepTrap attack. Whereas the formal procedure will be described in Section 3.4, we introduce the basic concepts in this subsection.

3.3.1 Object-User Graph

Figure 2 is created to represent the relationship between objects and users. On the left side of the graph, each vertex represents an object. On the right side of the graph, each vertex represents a user. Each object has a popularity score (a, b) , where a is the number of positive feedbacks and b is the number of negative feedbacks given to this object. If a user provides feedback to an object, a link is created between this user and the object. The link has weight 1, if the feedback from the user agrees with the quality of the object which is marked by the reputation system. Otherwise, the weight of the link is -1 . This graph is referred to as the *object-user graph*.

3.3.2 Correlation Between Objects

Based on the object-user graph, we can calculate the correlation between object O_i and O_j , denoted by $Cor(O_i, O_j)$. This correlation describes how O_j ’s reputation is affected if O_i becomes a trap. $Cor(O_i, O_j)$ is calculated in three steps.

1. Between O_i and O_j , find all the paths with length 2. For example, if user u_k provides feedbacks to both O_i and O_j , there exists a path $O_i - u_k - O_j$ and the path length is 2.
2. For each path, calculate the *path_effect* value as the product of the two links’ weights. Thus, the *path_effect* value can be either 1 or -1 . When *path_effect* = 1, the quality reputation of O_j will move opposite to its original value if O_i is turned into a trap. That is, high object reputation will become lower, and low object reputation will become higher. When *path_effect* = -1 , the object reputation will be maintained. That is, the effect of this path is to make high object reputation become higher and low object reputation become lower. The truth table (Table 2) supports the above statements, which will be explained later.
3. Let N_1 denote the number of paths with the positive *path_effect* values, and N_2 denote the number of paths

with negative *path_effect* values. Then, the correlation is calculated as

$$Cor(O_i, O_j) = N_1 - N_2. \quad (5)$$

It is noted that $Cor(O_i, O_j)$ can be used for trap selection. If $Cor(O_{i_1}, O_j) > Cor(O_{i_2}, O_j)$, turning O_{i_1} into a trap is more likely to have a greater impact on the reputation of O_j than turning O_{i_2} into a trap. In the example shown in Figure 2, $Cor(O_2, O_1) = 2$ and $Cor(O_5, O_1) = 1$.

To understand table 2, we examine the path $O_2 - u_1 - O_1$ shown in Figure 2 with the elements in the first row of Table 2 explained one by one, from left to right.

1. Initially, O_2 is marked as *high* quality by the system;
2. u_1 gives an honest feedback to O_2 which is *positive*;
3. Since u_1 's feedback agrees with O_2 's quality reputation, the weight of link $O_2 - u_1$ is 1;
4. If O_2 is turned into a trap, the system will mark O_2 as low quality and also think u_1 gives a wrong feedback, which leads to a *reduction* in u_1 's feedback reputation;
5. Initially, O_1 is marked as *high* quality by the system;
6. u_1 gives an honest feedback to O_1 which is *positive*;
7. Since u_1 's feedback agrees with O_1 's quality reputation, the weight of link $u_1 - O_1$ is 1;
8. Since the feedback reputation of u_1 , who gives a positive feedback, is reduced (see column 4), the object quality reputation of O_1 could also be *reduced*;
9. Making O_2 into a trap will make a high object quality reputation lower, that is, the object quality reputation of O_1 moves toward the *opposite* direction.
10. The *path_effect* value which is the product of two link weights, is 1.

The eight rows in Table 2 represent all different cases. It is seen that the *path_effect* value agrees with the direction of reputation change. It is a good indicator showing the effect of the traps.

3.3.3 Object-targeted RepTrap and User-targeted RepTrap

The RepTrap attack can mislead the object quality reputation as well as the feedback reputation of users.

- When the primary attack goal is to mislead the quality reputation of one or multiple specific objects, we call it the *object-targeted RepTrap*.
- When the primary attack goal is to mislead the feedback reputation of one or multiple specific users, we call it the *user-targeted RepTrap*.

Using the object-user graph, we can unify these two types of RepTrap attacks. In the user-targeted RepTrap, we create a *virtual object*. This virtual object is linked to all the users who are the target of the attacker. The weight of the links is 1. It is easy to see that reducing the feedback reputation of the users is equivalent to reducing the quality reputation of the virtual object. By creating the virtual object, we can convert the user-targeted RepTrap attack into the object-targeted RepTrap attack. Therefore, in the rest of the section 3, we only describe how to conduct the object-targeted RepTrap.

3.3.4 Knowledge Level of Attackers

In different application scenarios, the attackers may have different amount of knowledge about the reputation systems, which affects their ways to conduct the RepTrap attack. In this subsection, we discuss a few typical cases about the attackers' knowledge level.

The attacker wants to know all of the following information in the reputation system, including

- $info_1$: specific algorithms and parameters (especially thresholds) used in the reputation system
- $info_2$: the information in the object-user graph
- $info_3$: feedback reputation of malicious users
- $info_4$: feedback reputation of other users

First of all, we cannot assume that the algorithms used in the reputation systems are secrets. In other words, the security and robustness of the reputation system should not rely on the secrecy of the algorithms. This is a general principle in security research. Therefore, we assume that the attacker knows the algorithms in the reputation system.

Second, we assume that the attacker can judge whether an object has been marked as low-quality by the system. This assumption is reasonable because the system must take some actions to deal with low-quality objects, such as removing them from the system or ranking them extremely low compared to other normal objects. Based on this assumption, the attacker can insert positive and negative feedbacks to a newly published object, and examine how the reputation changes and when this object is detected as low-quality. By doing this for multiple times, the attacker can estimate the detection thresholds. Therefore, it leads to another assumption that the attacker can estimate the parameters in the algorithms, e.g. *threshold* in (3).

Third, in many reputation systems, the information about who provides what feedbacks to which objects is open to the public. In addition, the object quality information is obviously known to all users. In those systems such as Epinion, the object-user graph is known to the attacker. Of course, there are systems that do not publish details about the feedback information.

Furthermore, since the attacker knows the algorithms and the behavior of malicious users, the attacker surely knows the feedback reputation of malicious users. Some systems publish information about the trustworthiness of users, which is directly related to the users' feedback reputation. Examples include Epinion and Amazon. In some other systems, the attacker can easily dig out the history of users giving feedbacks. Knowing this history and the algorithms, the attacker can estimate the feedback reputation of good users. Of course, it is possible that the attackers cannot obtain this information in some systems.

Based on the above discussion, we define three knowledge levels:

- Knowledge Level 1 (KL1): the attacker knows $info_1$ and $info_3$.
- Knowledge Level 2 (KL2): the attacker knows $info_1$, $info_2$ and $info_3$.
- Knowledge Level 3 (KL3): the attacker knows $info_1$, $info_2$, $info_3$, and $info_4$.

In Section 3.4, we will first present the specific attack procedure for the KL3 case, and then extend the procedure to KL2 and KL1 cases.

$R^{qr}(O_i)$	u_k 's feedback to O_i	weight of link $O_i - u_k$	if O_i becomes trap, change in $R^{fr}(u_k)$	$R^{qr}(O_j)$	u_k 's feedback to O_j	weight of link $O_j - u_k$	change in $R^{qr}(O_j)$ due to change in $R^{fr}(u_k)$	maintain original $R^{qr}(O_j)$?	path effect value
high	positive	1	reduced	high	positive	1	reduced	opposite	1
high	positive	1	reduced	low	negative	1	increased	opposite	1
high	positive	1	reduced	high	negative	-1	increased	maintain	-1
high	positive	1	reduced	low	positive	-1	reduced	maintain	-1
high	negative	-1	increased	high	positive	1	increased	maintain	-1
high	negative	-1	increased	low	negative	1	reduced	maintain	-1
high	negative	-1	increased	high	negative	-1	reduced	opposite	1
high	negative	-1	increased	low	positive	-1	increased	opposite	1

Table 2: Truth table for verifying that *path_effect* value represents the effects of a trap upon O_j 's object reputation.

3.4 RepTrap Attack Procedure

Before we give a formal RepTrap Attack procedure, some important terminologies and notations are summarized.

- *Object Correlation* between object O_i and O_j is denoted by $Cor(O_i, O_j)$ (see Section 3.3.2).
- *Malicious Users* are the user IDs under the control of the attacker. We use m_i to denote a malicious user.
- *Attack Capability of a Malicious User* describes the impact of the feedback from this particular user. Let a_i denote the attack capability of the malicious user m_i . For example, we can choose a_i as the feedback reputation of m_i .
- *Attack Capability of a Set of Malicious User* describes the overall impact when every user in this set gives feedback to the same object. Let $A(S)$ denote the attack capability of user set S . The simplest way to calculate $A(S)$ is $A(S) = \sum_{i:m_i \in S} a_i$.
- *Hardness Value* of an object describes the difficulty level in turning this object into a trap. Let $H(O_k)$ denote the hardness of object O_k . When the object quality reputation is calculated using (2), one reasonable way to calculate $H(O_k)$ is

$$H(O_k) = R_{pos}^{fr}(O_k) \cdot \frac{1 - \text{threshold}}{\text{threshold}} - R_{neg}^{fr}(O_k), \quad (6)$$

where $R_{pos}^{fr}(O_k)$ (or $R_{neg}^{fr}(O_k)$) is the summation of feedback reputation of the users who have provided positive (or negative) feedbacks to O_k . In the case of KL3, the attacker can calculate $H(O_k)$ using (6).

- *Malicious User Set*, denoted by S_m , contains all the malicious users that are under the control of the attacker.
- *Target Object Set*, denoted by S_t , is the set of objects that are the attacker's target. That is, the attacker wants to reduce or boost the reputation of the objects in S_t .
- *Related Object Set*, denoted by S_r , contains the objects that have positive correlation with at least one object in S_t . In other words, turning an object in S_r into a trap will help to attack the objects in S_t .
- *Candidate Object Set*, denoted by S_c , contains the objects that satisfy two conditions:
 - They are marked as high-quality objects.
 - The objects' hardness values are smaller than the attack capability of the malicious user set. That is, if $O_k \in S_c$, then $H(O_k) < A(S_m)$.

In other words, the attacker only turns high-quality objects into traps, and the attacker does not try to create a trap unless the attacker has enough attack capability. Note that all traps are marked as low-quality. Since regular users usually do not use low-quality objects, traps will not receive many new feedbacks from honest users. On the other hand, if the system marks a low-quality object as high-quality, this object will be used by many users and quickly receives a large number of negative feedbacks. The system will correct its wrong judgement quickly.

- *Effectiveness Value* of an object describes how much influence the attacker can have on the target set if the attacker turns this particular object into a trap. The effectiveness value of the object O_k can be calculated as

$$E(O_k) = \sum_{j:O_j \in S_t} Cor(O_k, O_j) \quad (7)$$

- *Feedback Allocation* describes the specific behavior of the malicious users. From the attacker's point of view, *the most challenging job is to determine how many negative feedbacks should be given to which objects in order to create traps*, i.e. how to allocate the dishonest feedbacks. We introduce the notation

$$FA(S_c; S_{neg}^m(O_k) \text{ for } k : O_k \in S_c)$$

which means that the malicious users in set $S_{neg}^m(O_k)$ provide negative feedbacks to O_k , and O_k is in the candidate set S_c .

- *Round* is the time unit in which the attacker can adjust the feedback allocation. Without losing generality, one malicious user ID can provide up to one feedback in each round. The concept of round simplifies the description of the RepTrap attack procedure.

The core task of the RepTrap attack is to determine the feedback allocation in each round such that the quality reputation of the objects in the target set (S_t) is affected most. One round of the RepTrap attack is conducted in 9 steps.

In **Step 1**, construct the related object set (S_r) based on the definition given previously.

In **Step 2**, calculate the hardness value for each object in S_r , i.e. $H(O_j)$, for $j : O_j \in S_r$.

In **Step 3**, calculate the attack capability of each malicious user, as well as the malicious user set, i.e. $A(S_m)$.

In **Step 4**, construct the candidate object set (S_c) by comparing candidate $A(S_m)$ and the hardness values of the objects in S_r .

In **Step 5**, for each object in the candidate set (S_c), calculate its correlation with all other objects in the target object set (S_t).

In **Step 6**, calculate the effectiveness value for each object in S_c , i.e. $E(O_j)$, for $j : O_j \in S_c$.

In **Step 7**, determine the feedback allocation in this round by solving the following optimization problem.

inputs: $S_c; E(O_j)$ and $H(O_j)$ for $j : O_j \in S_c; a_i, m_i \in S_m$
 outputs: feedback allocation, i.e.
 $FA(S_c; S_{neg}^m(O_j))$ for $O_j \in S_c$

The optimization problem is:

$$\max_{FA(\cdot)} \sum_{j: S_{neg}^m(O_j) \neq \phi} E(O_j) \quad (8)$$

under constraints:

$$A(S_{neg}^m(O_j)) > H(O_j) \text{ if } S_{neg}^m(O_j) \neq \phi \quad (9)$$

$$S_{neg}^m(O_j) \cap S_{neg}^m(O_i) = \phi, \forall i, j : O_i \in S_c, O_j \in S_c, i \neq j \quad (10)$$

$$\bigcup_{j: O_j \in S_c} S_{neg}^m(O_j) \subseteq S_m \quad (11)$$

Here, (8) states that the attacker finds the feedback allocation, which maximizes the summation of the effectiveness values of the objects that will be turned into traps. (10) and (11) state that one malicious user can provide at most one feedback in one round. The constraint (9) is for ensuring the success of creating traps. That is, if the attacker provides negative feedbacks to an object, the attacker aims to turn the object into a trap in this round. In practice, we often overestimate the hardness value because it is possible that some good users provide honest feedbacks to the object in S_c during the process of the RepTrap attack. Based on the previous feedback history, the attacker can estimate the number of new feedbacks from honest users, and adjust the hardness value estimation accordingly.

In **Step 8**, the attacker inserts feedbacks according to the feedback allocation calculated in step 7.

In **Step 9**, at the end of each round, the attacker checks whether the traps are successfully created. If $O_i \in S_c$ is turned into a trap, set the trap indicator T_i as 1. (The default T_i value is 0). If O_i is not successfully turned into a trap for some reasons, O_i is put into the candidate set in the next round. This might happen only with a very small probability.

After one round the attack, the attacker can go back to step 1 and do more rounds, until the ending condition is met.

After reducing the good users' feedback reputation and enhancing their own feedback reputation through the RepTrap attack, the attacker is ready to attack the objects in the target set S_t . According to the specific attack goals, the attacker provides either positive or negative feedbacks to the objects in S_t .

In this 9-step procedure, two issues are not specified: the ending condition and the solution to the optimization problem in step 7. Next, we address these two issues.

The **ending condition** depends on the specific goal of attacker. Here are two examples.

- If the attacker's goal is to reduce or boost the quality reputation of the objects in S_t as much as possible, the ending condition is
 - the candidate set S_c is empty in a new round.

This ending condition means that no more traps can be created. In this case, the attacker creates as many traps as possible before giving dishonest feedbacks to the objects in S_t .

- If the attacker's goal is to make the system mark some high-quality object as low-quality, the ending condition is

- the candidate set S_c is empty;
- or the attacker estimates that he can achieve the attack goal given the users' feedback reputations in a new round.

If the attacker has knowledge KL3, he can easily perform this estimation. If the attacker has knowledge KL1 and KL2, he may try to attack the objects in S_t and see how the object reputation changes.

The **optimization** problem in step 7 is NP-hard. Therefore, we developed a sub-optimal solution as shown in Procedure 1.

Procedure 1 Feedback Allocation Algorithm

```

1: Set  $S_{neg}^m(O_k) = \phi$  for  $k : O_k \in S_c$ .
2: Let  $S_a$  denote the set of unused malicious users. Set  $S_a = S_m$ .
3: Let  $S_o$  denote the set of remaining objects that can be turned into traps. Set  $S_o = S_c$ .
4: Calculate the benefit value as  $B_k = E(O_k)/H(O_k)$ , for  $k : O_k \in S_c$ .
5: while  $A(S_o)$  is not empty do
6:   From  $S_o$ , select the object that has the largest benefit value. Assume that this object is  $O_r$ .
7:   From  $S_a$ , select a set of the malicious users, denoted by  $S_s$ , such that  $A(S_s) > H(O_r)$  and the number of malicious users in  $S_s$  is minimized.
8:   if the selection of  $S_s$  is successful then
9:     set  $S_{neg}^m(O_r) = S_s$ 
10:    remove  $S_s$  from  $S_a$ 
11:   end if
12:   Remove  $O_r$  from  $S_o$ 
13:   Remove the objects whose Hardness is not less than  $A(S_a)$  from  $S_o$ 
14: end while

```

To further enhance the power of RepTrap, we introduce the concept of *enhancement round*. In the enhancement round, the attacker finds an object O_e that is not in S_r . Then, the malicious users give honest feedback to O_e . This will not affect the quality reputation of the objects in the target set, but can increase the feedback reputation of malicious users, as well as their attack capability. Then the malicious users may turn more related objects into traps. In fact, the malicious users can also turn O_e into a trap, but it is easier to give honest feedback.

So after the ending condition of the regular RepTrap rounds is met, the attacker checks whether the hardness of the objects in the related candidate set S_r can be reached if the feedback reputation of the malicious users is increased. If not, there is no room for improvement. Otherwise, the attacker runs a few enhancement rounds to boost the feedback reputation of malicious users, which will improve their attack capability. Then, the attacker conducts regular RepTrap rounds until the ending condition is met again. The usage of the enhancement rounds can slightly improve the effectiveness of RepTrap.

3.5 RepTrap Attack for KL1 and KL2 Cases

The RepTrap attack procedure in Section 3.4 is obviously suitable for the case that the attacker has KL3. In this section, we briefly discuss how to conduct RepTrap for the KL1 and KL2 cases.

When the attacker does not know who provide feedbacks to the target set S_t (i.e. KL1), it is difficult to determine an effective strategy for creating traps. In this case, the attacker can simply try to create as many traps as possible. Note that creating traps not only make the objects in S_t easier to attack, but also make the system publish wrong object quality reputation.

When the attacker does not know the feedback reputation of good users, the attacker cannot calculate the hardness value $H(O_k)$ accurately. Instead, the attacker needs to estimate

$H(O_k)$. For example, the attacker can estimate $H(O_k)$ as

$$\hat{H}(O_k) = N(O_k) \cdot R_{avg}^{fr} \cdot \frac{1 - \text{threshold}}{\text{threshold}}, \quad (12)$$

where $N(O_k)$ is the number of users who have provided feedbacks to O_k (excluding the malicious users), and R_{avg}^{fr} is the average feedback reputation of good users. In most systems, $N(O_k)$ is known and the attacker only needs to estimate R_{avg}^{fr} . To estimate R_{avg}^{fr} , the attacker can use some additional user IDs that provide honest feedbacks to some other objects, just like good users. These additional users' feedback reputation can be used as a guideline for estimating R_{avg}^{fr} . Usually, R_{avg}^{fr} should be overestimated to ensure the success of creating traps.

4. PERFORMANCE EVALUATION

4.1 Experiment Description

As discussed in Section 1, we evaluate the performance of RepTrap attack in P2P file-sharing systems. We build a simulator using models similar to [11] and the parameters from Maze [18]. Maze is a popular P2P file-sharing system with more than 2 million registered users and 40,000 users online at peak time. In the simulation, there are 1,000 good users and 100 high quality files. Users follow Poisson process to download files and an important Zipf parameter, which is 0.676, is used to drive the downloading frequency. While another Zipf parameter, which is 0.890, is used to drive the popularity of the files [11]. Without loss of generality, we assume good users always provide honest feedbacks after successfully file downloads. (Of course, we can assume that good users provide honest feedbacks with a certain probability, but this will not have meaningful impact on the simulation results.) Each user can give no more than one feedback to a file. The reputation system model in the simulator is the same as that in Section 3.2.

We compare RepTrap with two other schemes: RepBad and RepSelf, which are described in Section 3.2. The attacks are performed in rounds and 50 malicious users are added. In one round of attack, each malicious user can only provide one feedback. With the RepBad strategy, malicious users attack the target files by providing false feedbacks in the first round. With the RepSelf strategy, malicious users give honest feedbacks to non-targeted files and gain feedback reputation. After getting enough feedback reputation, they provide false feedbacks to the target files. With the RepTrap strategy, malicious users create traps first and then attack the target files.

4.2 Attack model 1: Targeting a popular file

In a P2P file-sharing system, the popularity of files is driven by the Zipf parameter. More popular is a file, more feedbacks from good users are given to this file, and therefore more difficult for malicious users to conduct the attack. In this experiment, We attack the 1st (i.e. top 1) and the 10th (i.e. top 10) popular files and aim to make the system identify them as low quality files (i.e. faked files or pollution).

If the malicious users start an attack at the beginning of the simulation, the attack is much easier to succeed. This is because there are few feedbacks for each file and even a small number of false feedbacks can easily slander a good file. As the time goes, more feedbacks are provided to the files and the attack gets harder and harder to succeed with limited number of malicious users. We therefore change the starting time of the attack to show the attack performance measured by success probability. For each starting time, we conduct 1,000 simulations and success probability is calculated as the number of successful attacks divided by 1000.

Figure 3 and 4 are for the scenarios that malicious users attack the 1st and the 10th popular files respectively. The horizontal axis is attack starting time and the vertical axis is the




Figure 3: Success probability when attacking the 1st popular file




Figure 4: Success probability when attacking the 10th popular file

success probability. It is obvious that the proposed scheme has the best performance and RepBad has the worst performance. RepBad can only successfully attack the 1st popular file within 2 days and the 10th popular file within 10 days after the file is published. After that time, the success probability decreases sharply. RepSelf has better success probability than RepBad because in RepSelf malicious users can improve their feedback reputation and have more attack power. However, the performance of RepSelf is still much worse than RepTrap. For example, when RepSelf has only 10% chance of success, RepTrap has 60% and 80% of success probability in Figure 3 and 4 respectively. These results clearly show the advantage of reducing honest users' feedback reputation, which can only be achieved in RepTrap.

In the above experiment, both RepTrap and RepSelf need several attack rounds for the malicious users to gain feedback reputation. To compare the efficiency of RepTrap and RepSelf, we examine the *round number*, which is the number of rounds needed to gain enough feedback reputation to successfully attack a file. Please note that in each attack round, 50 malicious users provide 50 feedbacks (each user provides one feedback). The number of attack rounds therefore can represent the total amount of feedbacks provided by the malicious users.

Figure 5 and Figure 6 show the number of rounds needed




Figure 5: The number of rounds needed to successfully attack the 1st popular file




Figure 6: The number of rounds needed to successfully attack the 10th popular file




Figure 7: Maximum percentage of files turned into traps by 50 malicious users

to successfully attack the 1st and the 10th files, respectively. Three schemes are compared: RepSelf, RepTrap in KL3 case, and RepTrap in KL2 case. Recall that KL2 and KL3, defined in Section 3.3.4, represent different knowledge levels of the attacker. In KL2, the attacker has to estimate the good users' feedback reputation that affects the hardness value calculation. We assume it overestimates the hardness values of the objects by 20% in KL2. Compared with RepSelf, both RepTrap methods need much less number of rounds. This means that the RepTrap attackers spend much less resources. Compared with the KL3 case, the KL2 case needs slightly more rounds. It indicates that imperfect knowledge about the good users' feedback reputation will not largely reduce the power of the RepTrap attack.

4.3 Attack model 2: Blind attack

In this experiment, we consider the KL1 case, in which malicious users know neither the object-user graph nor the good users' feedback reputation. In this case, RepTrap does not target any specific files or users. Instead, it aims to disturb the entire reputation system by creating as many traps as possible. The malicious users will continue to make traps, starting from the least popular files to more popular files, until they cannot make any more traps.

Figure 7 shows the number of traps created by the 50 mali-




Figure 8: Reduction percentage of the average of feedback reputation for good users

cious users with different attack starting time. It is seen that at the beginning, all files can be turned into traps. The trap number decreases gradually with the starting time. However, even after day 50, more than 80% of the files can be trapped. This is because the popularity of the files in P2P file-sharing systems is driven by the power law distribution, so there are a large number of files downloaded by only a few users. From Figure 3, 4 and 7, we can conclude that RepTrap is a very powerful attack, which can attack most of the files no matter the starting time. The attack performance is only sensitive to the starting time when the attack targets very popular files.

Let R_{avg}^{wo} and R_{avg}^w denote the average of good users' feedback reputation without and with the RepTrap attack. In Figure 8, the vertical axis is *reduction percentage*, defined as $(R_{avg}^{wo} - R_{avg}^w)/R_{avg}^{wo}$. The horizontal axis is still the attack starting time. There observations are made.

First, at the very beginning, the reduction percentage is very low. This is because honest users have not given many feedbacks. The attacker cannot hurt a user's feedback reputation if this user does not provide feedbacks.

Second, as time goes on, more feedbacks are provided by honest users and the traps are more powerful to reduce honest users' feedback reputation. Therefore, we see that the reduction percentage gradually increases to 28% within 30 days.

Third, after day 30, the reduction percentage is not changing any more. This is because making traps becomes more difficult. Recall that it gets harder to attack the 10th popular file after day 30 (see Figure 4). Considering **there are 1,000 good users and only 50 malicious users** in the simulation, 28% reduction in average feedback reputation of all honest users is a big achievement. In practice, the attacker can control more user IDs as the time goes on, and therefore can even achieve better performance.

5. DISCUSSION

Through analysis and simulations, we have seen that RepTrap can severely degrade the quality reputation of objects and feedback reputation. Comparing RepBad and RepSelf, RepTrap needs the fewest user IDs to achieve the attack goal.

In fact, RepTrap has even broader influence. We use the P2P file-sharing system as an example to illustrate it.

First, RepTrap hurts the users' incentive to contribute. Nobody would like to see that the files they published are marked as faked files. The system with many traps would discourage people to publish high-quality files.

Second, many reputation systems calculate the reputation of the users in terms of publishing high quality files. This reputation can be referred to as the *user service reputation*. RepTrap obviously damage the user service reputation by making the system mark high-quality files as low-quality.

Third, many reputation systems have mechanisms to encourage good users and punish bad users. The target users of RepTrap would not receive the reward they deserve. Especially, the malicious users will not be punished. This will undermine the incentive for participation in the service or system.

Fourth, in this paper, we study the reputation systems using feedback reputation. If RepTrap is not solved, the system may have to give up using feedback reputation, which open doors to simple bad-mouthing attacks. Even if the system does not use feedback reputation, RepTrap still hurts the object quality reputation by slandering good files directly.

Therefore, a throughout understanding on this attack, as well as developing defense mechanisms, will have significant impact on the security and usability of reputation systems used today and in the future. This paper is the first effort toward understanding and solving this security problem in reputation systems. More research on defense is expected.

6. CONCLUSION

In this paper, we reported the discovery of a new and powerful attack, named RepTrap, against feedback-based reputation systems. Through an in-depth investigation, we presented the scenarios that this attack can be applied to, the ways to effectively conduct this attack, and the consequence of the attack. The performance evaluation is based on real data and realistic user models in a popular P2P file-sharing system. The simulation results as well as the case studies have demonstrated that the RepTrap attack can significantly reduce the resources required to attack popular objects. With the same resources, the success probability of attacks was greatly increased. The broader impact of this attack and several variations of RepTrap were also discussed. The investigation in this paper has moved the arms race between attack and defense in reputation systems to a new level.

7. REFERENCES

- [1] Emule specification. <http://www.emule.com>.
- [2] Kazaa. <http://www.kazaa.com>.
- [3] J. Brown and J. Morgan. Reputation in online auctions: The market for trust. *California Management Review*, 49(1):61–81, 2006.
- [4] D. Cosley, S. Lam, I. Albert, J. Konstan, and J. Riedl. Is seeing believing?: How recommender systems influence users' opinions. In *Proceedings of SIGCHI Conference on Human Factors in Computing Systems*, 2003.
- [5] C. Dellarocas. Immunizing online reputation reporting systems against unfair ratings and discriminatory behavior. In *Proceedings of ACM EC*, 2000.
- [6] C. Dellarocas. Mechanisms for coping with unfair ratings and discriminatory behavior in online reputation reporting systems. In *Proceedings of ICIS*, 2000.
- [7] C. Dellarocas. The digitization of word-of-mouth: Promise and challenges of online reputation systems. *Management Science*, 49(10):1407–1424, October 2003.
- [8] C. Dellarocas. Strategic manipulation of internet opinion forums: Implications for consumers and firms. *Management Science*, October 2006.
- [9] J. R. Douceur. The sybil attack. In *Proceedings of IPTPS*, March 2002.
- [10] Q. Feng and Y. Dai. Lip: A lifetime and popularity based ranking approach to filter out fake files in p2p file sharing systems. In *Proceedings of IPTPS*, 2007.
- [11] K. P. Gummadi, R. J. Dunn, S. Saroiu, S. D. Gribble, H. M. Levy, and J. Zahorjan. Measurement, modeling, and analysis of a peer-to-peer file-sharing workload. In *Proceedings of ACM SOSP*, 2003.
- [12] A. Harmon. Amazon glitch unmasks war of reviewers. *The New York Times*, February 14 2004.
- [13] M. Hines. Scammers gaming youtube ratings for profit. InfoWorld, http://www.infoworld.com/article/07/05/16/cybercrooks_gaming_google_1.html, May 2007.
- [14] K. Hoffman, D. Zage, and C. Nita-Rotaru. A survey of attack and defense techniques for reputation systems. Technical Report CSD TR #07-013, Purdue University, 2007.
- [15] A. Josang and R. Ismail. The beta reputation system. In *Proceedings of the 15th Bled Electronic Commerce Conference*, June 2002.
- [16] A. Josang, R. Ismail, and C. Boyd. A survey of trust and reputation systems for online service provision. *Decision Support System*, 43(2):618–644, 2007.
- [17] S. D. Kamvar, M. T. Schlosser, and H. Garcia-Molina. The eigentrust algorithm for reputation management in p2p networks. In *Proceedings of WWW*, May 2003.
- [18] Q. Lian, Z. Zhang, M. Yang, B. Y. Zhao, Y. Dai, and X. Li. An empirical study of collusion behavior in the maze p2p file-sharing system. In *Proceedings of ICDCS*, 2007.
- [19] S. Marti and H. Garcia-Molina. Taxonomy of trust: categorizing p2p reputation systems. *Computer Networks*, 50(4):472–484, 2006.
- [20] P. Resnick, R. Zeckhauser, E. Friedman, and K. Kuwabara. Reputation systems. *Communications of the ACM*, 43(12):45–48, 2000.
- [21] I. Stoica, R. Morris, D. Karger, F. Kaashoek, and H. Balakrishnan. Chord: A scalable Peer-To-Peer lookup service for internet applications. In *Proceedings of ACM SIGCOMM*, 2001.
- [22] Y. Sun, Z. Han, W. Yu, and K. J. R. Liu. A trust evaluation framework in distributed networks: Vulnerability analysis and defense against attacks. In *Proceedings of IEEE INFOCOM*, April 2006.
- [23] G. Swamynathan, B. Y. Zhao, and K. C. Almeroth. Decoupling service and feedback trust in a peer-to-peer reputation system. In *Proceedings of ISPA Workshops*, 2005.
- [24] K. Walsh and E. G. Sirer. Experience with an object reputation system for peer-to-peer filesharing. In *Proceedings of USENIX NSDI*, 2006.
- [25] J. Weng, C. Miao, and A. Goh. An entropy-based approach to protecting rating systems from unfair testimonies. *IEICE TRANSACTIONS on Information and Systems*, E89-D(9):2502–2511, September 2006.
- [26] M. Yang, Q. Feng, Y. Dai, and Z. Zhang. A multi-dimensional reputation system combined with trust and incentive mechanisms in p2p file sharing systems. In *Proceedings of ICDCS Workshops*, 2007.
- [27] Y. Yang, Y. Sun, J. Ren, and Q. Yang. Building trust in online rating systems through signal modeling. In *Proceedings of ICDCS Workshops*, 2007.
- [28] H. Yu, P. Gibbons, M. Kaminsky, and F. Xiao. Sybillimit: A near-optimal social network defense against sybil attacks. In *Proceedings of IEEE Symposium on Security and Privacy*, May 2008.
- [29] H. Yu, M. Kaminsky, P. Gibbons, and A. Flaxman. Sybilguard: Defending against sybil attacks via social networks. In *Proceedings of ACM SIGCOMM*, September 2006.